

Kafka Usecase

Gabriel D Hofer

July 22, 2021

Using single-node multi-broker setup , do the following: Create a topic called “logs” with 3 partitions and a replication factor of 3.

```
1 $ bin/kafka-topics.sh \  
2 --zookeeper localhost:2181 \  
3 --create --topic logs \  
4 --partitions 3 \  
5 --replication-factor 3
```

Describe the topic to understand in broker are the leader and in-sync replicas of your topic.

```
1 $ bin/kafka-topics.sh --zookeeper localhost:2181 --describe --topic logs
2 Topic: logs TopicId: Y7fJeVCRzWNGLWHAPE_bw PartitionCount: 3 ReplicationFactor: 3   Configs:
3   Topic: logs Partition: 0   Leader: 3 Replicas: 3,2,1 Isr: 3,2,1
4   Topic: logs Partition: 1   Leader: 1 Replicas: 1,3,2 Isr: 1,3,2
5   Topic: logs Partition: 2   Leader: 2 Replicas: 2,1,3 Isr: 2,1,3
```

Produce some log messages to the “logs” topic as key-value pairs where log-level is the key and log message is the value. Ex: A message can look like “WARN: Message not properly formatted”. Use Kafka console producer for doing this.

[illegible]

Launch three instances of Kafka Console Consumers as a consumer group and demonstrate how messages are distributed across three partitions and how each consumer in the group is consuming from each of the three partitions.

Consumer instance one:

```
1 $ ./bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic logs --group
  logs-group
2 ERROR: php 21/05/2015
3 ERROR: RailsApp 05/08/2015
4 ERROR: php 13/02/2016
5 DEBUG: php 22/08/2015
6 DONE: mysql 28/06/2015
7 DONE: php 11/01/2016
8 ERROR: php 22/11/2014
9 WARN: mysql 28/06/2015
```

Consumer instance two:

```
1 $ ./bin/kafka-console-consumer.sh --bootstrap-server localhost:9093 --topic logs --group
  logs-group
2 DEBUG: php 11/01/2016
3 ERROR: php 19/05/2015
4 ERROR: RailsApp 25/12/2015
5 ERROR: RailsApp 06/05/2015
6 INFO: php 19/05/2015
7 INFO: php 13/02/2016
8 ERROR: mysql 18/03/2015
9 ERROR: RailsApp 06/05/2015
```

Consumer instance three:

```
1 $ ./bin/kafka-console-consumer.sh --bootstrap-server localhost:9094 --topic logs --group
  logs-group
2 DONE: mysql 23/01/2016
3 ERROR: php 22/11/2014
4 ERROR: mysql 18/03/2015
5 ERROR: php 09/03/2015
6 ERROR: php 21/05/2015
7 WARN: RailsApp 05/08/2015
8 WARN: mysql 23/01/2016
9 ERROR: RailsApp 25/12/2015
10 WARN: php 22/08/2015
11 ERROR: php 09/03/2015
```

Write a simple Python client application using Kafka Producer API to producer messages to the “logs” topic. Using Fire-and-Forget approach for fastest message delivery. Send only “ERROR” messages.

```
1 from kafka import KafkaProducer
2 from time import sleep
3 import json
4 from datetime import datetime
5 import random
6
7 producer = KafkaProducer(bootstrap_servers=['localhost:9092'])
8 producer.send('logs', b'Client application using Kafka Producer API')
9
10 error_messages = [
11     '500 Internal Server Error',
12     '501 Not Implemented',
13     '502 Bad Gateway',
14     '503 Service Unavailable',
15     '504 Gateway Timeout',
16     '505 HTTP Version Not Supported',
17     '506 Variant Also Negotiates (RFC 2295)',
18     '507 Insufficient Storage (WebDAV; RFC 4918)',
19     '508 Loop Detected (WebDAV; RFC 5842)',
20     '510 Not Extended (RFC 2774)',
21     '511 Network Authentication Required (RFC 6585)',
22 ]
23
24 for i in range(50):
25     message = " ERROR : {} ".format(str( random.choice(error_messages) ))
26     producer.send('logs', json.dumps(message).encode('utf-8'))
27     sleep(0.5)
28     print("logs message ",i)
```

Write a simple Python client application using Kafka Consumer API to consume the data from the “logs” topic. Print the log messages on the Console. Turn-off auto-commit and explicit commit the messages from the application.

```
1 from kafka import KafkaConsumer
2 consumer = KafkaConsumer('logs', bootstrap_servers=['localhost:9092'])
3
4 # Turn off auto-commit messages
5 KafkaConsumer(enable_auto_commit=False, auto_commit_interval_ms=False)
6
7 for message in consumer:
8     print(message.value)
```